


Controlled q-label heterogeneity: iid | APPLE-Qinit diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	67.9% ± 1.6%	72.3% ± 1.5%	<u>90.0% ± 1.5%</u>	<u>85.0% ± 1.7%</u>	67.5% ± 3.5%	<u>66.8% ± 2.2%</u>	50.9% ± 1.2%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-Qinit	<u>69.5% ± 0.0%</u>	<u>72.6% ± 0.0%</u>	88.6% ± 0.4%	81.9% ± 1.5%	<u>69.4% ± 1.9%</u>	64.0% ± 1.7%	<u>56.8% ± 1.4%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-QInit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>70.1% ± 3.2%</u>	<u>74.3% ± 2.5%</u>	<u>93.1% ± 1.1%</u>	<u>86.7% ± 1.6%</u>	<u>70.3% ± 4.6%</u>	<u>67.5% ± 2.2%</u>	53.8% ± 5.0%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-QInit	69.2% ± 4.1%	72.4% ± 4.3%	88.9% ± 3.2%	81.5% ± 8.4%	69.6% ± 6.6%	65.9% ± 4.9%	<u>55.9% ± 3.3%</u>


Validation loss: baselines vs selected APPLE-QInit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-QInit micro / macro F1 (after-local selected model)



Pooled APPLE-QInit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-QInit mismatch gap




family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-QInit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>69.5% ± 0.8%</u>	<u>73.5% ± 0.6%</u>	<u>90.4% ± 0.9%</u>	<u>85.2% ± 1.6%</u>	<u>67.7% ± 0.6%</u>	<u>71.2% ± 1.1%</u>	52.9% ± 2.1%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-QInit	68.3% ± 5.0%	71.6% ± 3.9%	88.9% ± 1.5%	83.3% ± 1.2%	65.9% ± 9.0%	62.6% ± 8.4%	<u>57.5% ± 2.3%</u>


Validation loss: baselines vs selected APPLE-QInit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-QInit micro / macro F1 (after-local selected model)



Pooled APPLE-QInit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-QInit mismatch gap




family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	64.8% ± 3.3%	70.1% ± 2.7%	<u>92.2% ± 3.4%</u>	82.7% ± 3.5%	65.7% ± 8.3%	60.9% ± 7.0%	48.9% ± 2.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-Qinit	<u>69.0% ± 3.0%</u>	<u>72.6% ± 3.0%</u>	87.8% ± 5.3%	<u>86.0% ± 3.2%</u>	<u>72.2% ± 3.0%</u>	<u>61.1% ± 7.0%</u>	<u>55.8% ± 1.1%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	70.5% ± 1.9%	74.2% ± 2.2%	87.8% ± 6.8%	<u>86.4% ± 2.4%</u>	<u>74.0% ± 2.0%</u>	<u>69.7% ± 1.8%</u>	53.3% ± 1.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-Qinit	<u>72.6% ± 1.5%</u>	<u>75.1% ± 1.5%</u>	<u>89.3% ± 1.3%</u>	83.5% ± 4.4%	73.4% ± 2.4%	69.0% ± 1.8%	<u>60.2% ± 2.4%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	64.5% ± 3.5%	69.6% ± 3.3%	87.3% ± 2.7%	<u>84.4% ± 3.4%</u>	62.5% ± 7.2%	<u>67.0% ± 5.0%</u>	46.8% ± 1.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-Qinit	<u>68.5% ± 5.5%</u>	<u>71.4% ± 4.8%</u>	<u>88.8% ± 1.5%</u>	78.0% ± 8.3%	<u>69.2% ± 4.7%</u>	65.5% ± 6.4%	<u>55.7% ± 5.6%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




APPLE-Qinit training-loss decomposition | lid heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000000




C3-specialist | graph_id=3000001




C4-specialist | graph_id=3000002



C5-specialist | graph_id=3000003








C6-specialist | graph_id=3000004



APPLE-QInit task-wise DR matrices: receiver clients × donor task-heads | iid


family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000




APPLE-QInit task-wise DR dynamics | iid

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task




Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment




Controlled q-label heterogeneity | Mid | APPLE-Qinit diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	61.6% ± 4.0%	66.0% ± 4.3%	81.6% ± 7.0%	<u>78.2% ± 5.1%</u>	64.0% ± 3.6%	58.6% ± 5.2%	44.4% ± 3.0%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>68.0% ± 4.2%</u>	<u>70.8% ± 4.4%</u>	<u>85.2% ± 5.4%</u>	77.7% ± 7.8%	<u>68.7% ± 3.8%</u>	<u>66.5% ± 1.9%</u>	<u>54.7% ± 3.9%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-Qinit diagnostics | family=q_task_cycle2 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	63.8% ± 5.4%	68.0% ± 5.8%	<u>85.6% ± 8.0%</u>	80.6% ± 4.3%	65.4% ± 6.2%	59.3% ± 9.8%	49.2% ± 2.7%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>69.4% ± 5.3%</u>	<u>72.0% ± 6.1%</u>	85.2% ± 8.8%	<u>81.6% ± 11.2%</u>	<u>68.6% ± 7.6%</u>	<u>67.0% ± 1.0%</u>	<u>57.5% ± 2.5%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-Qinit diagnostics | family=q_task_cycle3 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	60.1% ± 11.6%	64.2% ± 10.9%	75.3% ± 14.5%	<u>77.9% ± 12.0%</u>	63.8% ± 6.1%	60.7% ± 11.5%	43.1% ± 11.1%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>67.7% ± 4.4%</u>	<u>70.4% ± 3.9%</u>	<u>85.7% ± 0.8%</u>	76.5% ± 7.3%	<u>69.6% ± 6.9%</u>	<u>65.2% ± 3.0%</u>	<u>55.0% ± 4.5%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	58.7% ± 5.5%	63.8% ± 5.1%	82.7% ± 2.7%	76.5% ± 5.8%	63.3% ± 5.5%	54.2% ± 7.8%	42.1% ± 3.8%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>68.0% ± 4.0%</u>	<u>71.1% ± 4.0%</u>	<u>89.1% ± 4.0%</u>	<u>79.1% ± 7.4%</u>	<u>64.8% ± 2.9%</u>	<u>66.6% ± 2.0%</u>	<u>55.8% ± 4.3%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-Qinit diagnostics | family=q_task_cycle5 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	61.2% ± 0.9%	66.3% ± 1.4%	<u>86.3% ± 4.7%</u>	<u>79.8% ± 5.8%</u>	63.6% ± 1.5%	58.1% ± 1.4%	43.8% ± 3.4%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>66.0% ± 6.4%</u>	<u>68.7% ± 6.5%</u>	80.9% ± 10.2%	75.0% ± 10.6%	<u>70.7% ± 6.2%</u>	<u>64.8% ± 3.5%</u>	<u>52.2% ± 6.5%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-Qinit diagnostics | family=q_task_cycle6 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	64.1% ± 1.8%	67.8% ± 2.4%	81.8% ± 8.2%	<u>78.3% ± 6.8%</u>	64.9% ± 1.4%	64.9% ± 4.4%	49.0% ± 3.2%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-Qinit	<u>68.7% ± 2.7%</u>	<u>71.8% ± 2.8%</u>	<u>86.3% ± 5.1%</u>	78.0% ± 7.6%	<u>71.8% ± 4.2%</u>	<u>69.1% ± 0.7%</u>	<u>54.0% ± 3.2%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




APPLE-QInit training-loss decomposition | Mid heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000005




C3-specialist | graph_id=3000006




C4-specialist | graph_id=3000007



C5-specialist | graph_id=3000008








C6-specialist | graph_id=3000009



APPLE-QInit task-wise DR matrices: receiver clients × donor task-heads | Mid


family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120




APPLE-QInit task-wise DR dynamics | Mid

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
 heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task




Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment




Controlled q-label heterogeneity: High | APPLE-Qinit diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	27.6% ± 1.8%	27.6% ± 1.8%	22.4% ± 1.5%	27.5% ± 2.1%	27.3% ± 1.2%	26.7% ± 1.2%	30.2% ± 1.5%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-Qinit	<u>64.4% ± 6.6%</u>	<u>66.6% ± 7.5%</u>	<u>73.3% ± 14.1%</u>	<u>76.2% ± 9.9%</u>	<u>64.7% ± 4.4%</u>	<u>62.5% ± 8.2%</u>	<u>53.2% ± 4.4%</u>


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-QInit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	28.6% ± 0.7%	28.6% ± 0.7%	23.4% ± 1.7%	29.1% ± 1.2%	30.2% ± 0.8%	28.5% ± 1.4%	32.0% ± 1.3%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-QInit	61.4% ± 6.7%	63.8% ± 7.0%	70.1% ± 11.0%	75.2% ± 13.6%	62.9% ± 1.5%	59.6% ± 10.5%	51.2% ± 5.9%


Validation loss: baselines vs selected APPLE-QInit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-QInit micro / macro F1 (after-local selected model)



Pooled APPLE-QInit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-QInit mismatch gap




family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	26.5% ± 4.7%	26.6% ± 4.7%	22.3% ± 4.2%	27.0% ± 7.0%	26.9% ± 5.3%	27.0% ± 4.1%	29.5% ± 3.2%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-Qinit	66.4% ± 5.6%	68.1% ± 5.9%	75.6% ± 8.9%	73.4% ± 4.4%	69.9% ± 5.2%	66.4% ± 9.2%	55.0% ± 5.0%


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




Controlled q-label heterogeneity: High | APPLE-QInit diagnostics | family=q_task_cycle4 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3% heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-QInit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	27.5% ± 1.5%	27.5% ± 1.4%	23.1% ± 1.8%	28.2% ± 0.7%	26.4% ± 1.8%	28.3% ± 1.7%	31.5% ± 2.1%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-QInit	63.5% ± 11.6%	65.7% ± 13.0%	74.1% ± 21.6%	77.4% ± 15.5%	60.9% ± 10.9%	63.5% ± 12.7%	52.8% ± 7.6%


Validation loss: baselines vs selected APPLE-QInit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-QInit micro / macro F1 (after-local selected model)



Pooled APPLE-QInit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-QInit mismatch gap




family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	24.2% ± 1.1%	24.2% ± 1.1%	20.1% ± 0.7%	25.3% ± 1.1%	24.5% ± 1.6%	22.9% ± 1.3%	28.3% ± 1.1%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-Qinit	64.1% ± 8.0%	66.2% ± 8.9%	74.4% ± 17.2%	77.0% ± 10.2%	67.3% ± 5.0%	57.5% ± 9.8%	55.0% ± 4.3%


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-Qinit, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	31.2% ± 8.8%	31.1% ± 8.9%	26.7% ± 8.1%	32.1% ± 9.5%	33.6% ± 11.2%	31.3% ± 9.2%	32.1% ± 6.8%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-Qinit	66.7% ± 4.1%	69.1% ± 4.6%	74.9% ± 10.9%	80.2% ± 6.5%	65.4% ± 3.3%	71.7% ± 2.3%	53.5% ± 2.6%


Validation loss: baselines vs selected APPLE-Qinit and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-Qinit micro / macro F1 (after-local selected model)



Pooled APPLE-Qinit micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-Qinit mismatch gap




APPLE-Qinit training-loss decomposition | High heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000010




C3-specialist | graph_id=3000011




C4-specialist | graph_id=3000012



C5-specialist | graph_id=3000013




C6-specialist | graph_id=3000014




APPLE-QInit task-wise DR matrices: receiver clients × donor task-heads | High

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


DR matrix for task cycle2: at best selected checkpoint




DR matrix for task cycle3: at best selected checkpoint




DR matrix for task cycle4: at best selected checkpoint



DR matrix for task cycle5: at best selected checkpoint




DR matrix for task cycle6: at best selected checkpoint




APPLE-QInit task-wise DR dynamics | High

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
 heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task



Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment

