


Controlled q-label heterogeneity: iid | APPLE-TaskHead diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>67.9% ± 1.6%</u>	<u>72.3% ± 1.5%</u>	<u>90.0% ± 1.5%</u>	<u>85.0% ± 1.7%</u>	<u>67.5% ± 3.5%</u>	<u>66.8% ± 2.2%</u>	50.9% ± 1.2%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	67.4% ± 3.8%	70.0% ± 3.5%	85.0% ± 2.0%	75.6% ± 4.1%	67.3% ± 5.2%	65.0% ± 5.1%	<u>54.8% ± 3.1%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>70.1% ± 3.2%</u>	<u>74.3% ± 2.5%</u>	<u>93.1% ± 1.1%</u>	<u>86.7% ± 1.6%</u>	<u>70.3% ± 4.6%</u>	<u>67.5% ± 2.2%</u>	<u>53.8% ± 5.0%</u>
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	63.7% ± 5.0%	66.7% ± 5.2%	84.8% ± 3.6%	71.8% ± 10.1%	65.5% ± 6.7%	59.8% ± 3.6%	51.9% ± 4.4%


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>69.5% ± 0.8%</u>	<u>73.5% ± 0.6%</u>	<u>90.4% ± 0.9%</u>	<u>85.2% ± 1.6%</u>	<u>67.7% ± 0.6%</u>	<u>71.2% ± 1.1%</u>	52.9% ± 2.1%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	66.7% ± 5.8%	69.6% ± 5.5%	84.3% ± 5.8%	81.0% ± 3.7%	65.1% ± 7.6%	64.5% ± 8.0%	<u>53.2% ± 3.8%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	64.8% ± 3.3%	70.1% ± 2.7%	<u>92.2% ± 3.4%</u>	<u>82.7% ± 3.5%</u>	65.7% ± 8.3%	60.9% ± 7.0%	48.9% ± 2.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	<u>69.5% ± 0.6%</u>	<u>71.5% ± 1.4%</u>	82.4% ± 5.7%	79.0% ± 3.8%	<u>71.0% ± 3.6%</u>	<u>67.5% ± 2.3%</u>	<u>57.7% ± 1.0%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>70.5% ± 1.9%</u>	<u>74.2% ± 2.2%</u>	<u>87.8% ± 6.8%</u>	<u>86.4% ± 2.4%</u>	<u>74.0% ± 2.0%</u>	<u>69.7% ± 1.8%</u>	53.3% ± 1.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	65.9% ± 7.2%	68.5% ± 6.6%	87.1% ± 1.9%	70.1% ± 9.7%	66.8% ± 7.8%	63.8% ± 9.5%	<u>54.8% ± 5.6%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	64.5% ± 3.5%	69.6% ± 3.3%	87.3% ± 2.7%	<u>84.4% ± 3.4%</u>	62.5% ± 7.2%	67.0% ± 5.0%	46.8% ± 1.7%
Local-Centralized	85.0% ± 0.7%	87.0% ± 0.6%	98.0% ± 1.1%	94.5% ± 1.4%	89.7% ± 1.8%	80.3% ± 0.5%	72.5% ± 1.3%
APPLE-TaskHead	<u>71.0% ± 2.8%</u>	<u>73.4% ± 2.6%</u>	<u>87.4% ± 5.8%</u>	80.2% ± 5.6%	<u>70.0% ± 3.2%</u>	<u>71.8% ± 2.0%</u>	<u>57.7% ± 2.3%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




APPLE-TaskHead training-loss decomposition | lid heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000000




C3-specialist | graph_id=3000001




C4-specialist | graph_id=3000002



C5-specialist | graph_id=3000003








C6-specialist | graph_id=3000004



APPLE-TaskHead task-wise DR matrices: receiver clients × donor task-heads | lid


family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000




APPLE-TaskHead task-cycle DR dynamics | iid

family=q_task_cycle2 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
 heterogeneity=iid | gamma=0.20 | task_jsd_mean=0.000 | task_jsd_median=0.000 | task_jsd_max=0.000 | family_centroid_jsd_mean=0.000


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task




Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment




Controlled q-label heterogeneity: mid | APPLE-TaskHead diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>61.6% ± 4.0%</u>	<u>66.0% ± 4.3%</u>	<u>81.6% ± 7.0%</u>	<u>78.2% ± 5.1%</u>	<u>64.0% ± 3.6%</u>	<u>58.6% ± 5.2%</u>	44.4% ± 3.0%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	58.9% ± 1.7%	60.8% ± 1.6%	72.0% ± 4.9%	62.1% ± 2.1%	60.7% ± 2.3%	55.9% ± 3.2%	<u>50.4% ± 2.3%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>63.8% ± 5.4%</u>	<u>68.0% ± 5.8%</u>	<u>85.6% ± 8.0%</u>	<u>80.6% ± 4.3%</u>	<u>65.4% ± 6.2%</u>	<u>59.3% ± 9.8%</u>	49.2% ± 2.7%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	58.0% ± 1.5%	60.3% ± 1.5%	75.7% ± 6.0%	64.2% ± 1.7%	59.0% ± 3.9%	52.7% ± 4.0%	<u>49.8% ± 1.7%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-TaskHead diagnostics | family=q_task_cycle3 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	60.1% ± 11.6%	64.2% ± 10.9%	75.3% ± 14.5%	<u>77.9% ± 12.0%</u>	63.8% ± 6.1%	<u>60.7% ± 11.5%</u>	43.1% ± 11.1%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	<u>64.7% ± 3.1%</u>	<u>67.5% ± 3.2%</u>	<u>87.1% ± 7.4%</u>	68.6% ± 5.5%	<u>67.6% ± 6.9%</u>	59.8% ± 6.1%	<u>54.3% ± 3.3%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>58.7% ± 5.5%</u>	<u>63.8% ± 5.1%</u>	<u>82.7% ± 2.7%</u>	<u>76.5% ± 5.8%</u>	<u>63.3% ± 5.5%</u>	54.2% ± 7.8%	42.1% ± 3.8%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	57.1% ± 3.7%	58.8% ± 3.3%	69.4% ± 7.5%	62.3% ± 2.3%	54.7% ± 5.0%	<u>57.9% ± 7.0%</u>	<u>49.6% ± 4.5%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-TaskHead diagnostics | family=q_task_cycle5 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>61.2% ± 0.9%</u>	<u>66.3% ± 1.4%</u>	<u>86.3% ± 4.7%</u>	<u>79.8% ± 5.8%</u>	63.6% ± 1.5%	<u>58.1% ± 1.4%</u>	43.8% ± 3.4%
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	58.5% ± 2.8%	59.5% ± 2.9%	66.7% ± 6.7%	57.9% ± 3.7%	<u>65.3% ± 1.6%</u>	54.8% ± 3.1%	<u>53.1% ± 3.6%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: Mid | APPLE-TaskHead diagnostics | family=q_task_cycle6 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3% | heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	<u>64.1% ± 1.8%</u>	<u>67.8% ± 2.4%</u>	<u>81.8% ± 8.2%</u>	<u>78.3% ± 6.8%</u>	<u>64.9% ± 1.4%</u>	<u>64.9% ± 4.4%</u>	<u>49.0% ± 3.2%</u>
Local-Centralized	84.3% ± 1.6%	86.3% ± 1.4%	96.7% ± 1.1%	94.7% ± 0.6%	87.4% ± 4.2%	80.9% ± 2.2%	71.8% ± 3.9%
APPLE-TaskHead	56.3% ± 1.0%	58.1% ± 2.0%	67.6% ± 10.3%	59.2% ± 1.6%	60.4% ± 0.8%	56.5% ± 3.8%	47.0% ± 0.5%


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




APPLE-TaskHead training-loss decomposition | Mid heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000005




C3-specialist | graph_id=3000006




C4-specialist | graph_id=3000007



C5-specialist | graph_id=3000008








C6-specialist | graph_id=3000009



APPLE-TaskHead task-wise DR matrices: receiver clients x donor task-heads | Mid


family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120




APPLE-TaskHead task-wise DR dynamics | Mid

family=q_task_cycle2 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 0.8%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.7%±0.0% / true 5.8% | cycle3=visible 0.8%±0.0% / true 6.2% | cycle4=visible 0.8%±0.0% / true 6.6% | cycle5=visible 0.8%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
 heterogeneity=mid | gamma=0.50 | task_jsd_mean=0.120 | task_jsd_median=0.121 | task_jsd_max=0.123 | family_centroid_jsd_mean=0.120


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task




Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment




Controlled q-label heterogeneity: High | APPLE-TaskHead diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	27.6% ± 1.8%	27.6% ± 1.8%	22.4% ± 1.5%	27.5% ± 2.1%	27.3% ± 1.2%	26.7% ± 1.2%	30.2% ± 1.5%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	31.1% ± 2.6%	31.1% ± 2.6%	25.6% ± 1.9%	31.0% ± 3.0%	31.4% ± 3.3%	30.7% ± 2.5%	34.1% ± 2.9%


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	28.6% ± 0.7%	28.6% ± 0.7%	23.4% ± 1.7%	29.1% ± 1.2%	30.2% ± 0.8%	28.5% ± 1.4%	32.0% ± 1.3%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	<u>33.0% ± 5.1%</u>	<u>33.1% ± 5.1%</u>	<u>27.1% ± 4.7%</u>	<u>34.0% ± 5.2%</u>	<u>34.7% ± 6.8%</u>	<u>33.5% ± 5.1%</u>	<u>36.0% ± 3.7%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: High | APPLE-TaskHead diagnostics | family=q_task_cycle3 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	26.5% ± 4.7%	26.6% ± 4.7%	22.3% ± 4.2%	27.0% ± 7.0%	26.9% ± 5.3%	27.0% ± 4.1%	29.5% ± 3.2%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	<u>33.6% ± 0.3%</u>	<u>33.6% ± 0.4%</u>	<u>29.3% ± 1.4%</u>	<u>32.6% ± 1.5%</u>	<u>33.9% ± 0.5%</u>	<u>35.1% ± 1.4%</u>	<u>37.0% ± 0.9%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: High | APPLE-TaskHead diagnostics | family=q_task_cycle4 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	27.5% ± 1.5%	27.5% ± 1.4%	23.1% ± 1.8%	28.2% ± 0.7%	26.4% ± 1.8%	28.3% ± 1.7%	31.5% ± 2.1%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	30.4% ± 3.2%	30.4% ± 3.2%	26.1% ± 2.7%	30.8% ± 3.4%	30.0% ± 3.9%	30.9% ± 3.5%	33.9% ± 3.3%


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: High | APPLE-TaskHead diagnostics | family=q_task_cycle5 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	24.2% ± 1.1%	24.2% ± 1.1%	20.1% ± 0.7%	25.3% ± 1.1%	24.5% ± 1.6%	22.9% ± 1.3%	28.3% ± 1.1%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	26.6% ± 1.9%	26.5% ± 2.0%	22.2% ± 1.2%	27.3% ± 2.5%	27.2% ± 2.4%	25.2% ± 1.6%	30.8% ± 2.3%


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)



Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




Controlled q-label heterogeneity: High | APPLE-TaskHead diagnostics | family=q_task_cycle6 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64_headmulti

family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3% | heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Best test performance (Fully Local vs Local-Centralized vs APPLE-TaskHead, rounds=80)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Fully Local	31.2% ± 8.8%	31.1% ± 8.9%	<u>26.7% ± 8.1%</u>	<u>32.1% ± 9.5%</u>	<u>33.6% ± 11.2%</u>	31.3% ± 9.2%	32.1% ± 6.8%
Local-Centralized	84.9% ± 1.3%	86.7% ± 0.8%	96.8% ± 1.6%	95.1% ± 1.4%	86.6% ± 2.5%	81.4% ± 1.3%	73.7% ± 5.8%
APPLE-TaskHead	<u>31.9% ± 4.7%</u>	<u>31.8% ± 4.7%</u>	26.1% ± 3.9%	31.9% ± 4.6%	<u>33.6% ± 5.4%</u>	<u>32.8% ± 4.9%</u>	<u>34.7% ± 4.9%</u>


Validation loss: baselines vs selected APPLE-TaskHead and end-of-round mismatch (rounds=80)




Pooled baseline micro / macro F1




Pooled APPLE-TaskHead micro / macro F1 (after-local selected model)



Pooled APPLE-TaskHead micro / macro F1 (end-of-round diagnostic only)




Delta diagnostics: baseline gaps and APPLE-TaskHead mismatch gap




APPLE-TaskHead training-loss decomposition | High heterogeneity

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


APPLE lambda scheduler | apple_mu=0.001




C2-specialist | graph_id=3000010




C3-specialist | graph_id=3000011




C4-specialist | graph_id=3000012



C5-specialist | graph_id=3000013




C6-specialist | graph_id=3000014




APPLE-TaskHead task-wise DR matrices: receiver clients × donor task-heads | High

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


DR matrix for task cycle2: at best selected checkpoint




DR matrix for task cycle3: at best selected checkpoint




DR matrix for task cycle4: at best selected checkpoint



DR matrix for task cycle5: at best selected checkpoint




DR matrix for task cycle6: at best selected checkpoint




APPLE-TaskHead task-wise DR dynamics | High

family=q_task_cycle2 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle3 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle4 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle5 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 0.3%±0.0% / true 6.3%
 family=q_task_cycle6 | n_clients=1 | cycle2=visible 0.3%±0.0% / true 5.8% | cycle3=visible 0.3%±0.0% / true 6.2% | cycle4=visible 0.3%±0.0% / true 6.6% | cycle5=visible 0.3%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
 heterogeneity=high | gamma=0.80 | task_jsd_mean=0.399 | task_jsd_median=0.399 | task_jsd_max=0.403 | family_centroid_jsd_mean=0.399


Task-head DR self-weight by task




Task-head DR row sum by task




Task-head off-diagonal absolute mass by task



Task-head DR movement away from p0 by task



Task-head off-diagonal negative mass by task



Task-head DR alignment

