


Controlled task-specialization level: Low | FedAvg diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64

family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
family=specialized_cycle3 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
family=specialized_cycle4 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
family=specialized_cycle5 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
family=specialized_cycle6 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3%
specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	83.4% ± 0.3%	85.7% ± 0.3%	96.8% ± 0.6%	94.9% ± 0.4%	86.8% ± 0.2%	78.7% ± 0.7%	70.4% ± 0.6%
FedAvg	81.6% ± 0.3%	84.0% ± 0.3%	96.2% ± 0.8%	95.4% ± 0.7%	82.3% ± 2.8%	78.0% ± 0.2%	68.4% ± 2.8%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




Controlled task-specialization level: Low | FedAvg diagnostics | family=specialized_cycle2 | n_clients=1 | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64

family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3% | specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.5% ± 1.5%	84.0% ± 1.5%	96.6% ± 1.3%	93.0% ± 1.5%	85.5% ± 2.0%	78.8% ± 2.1%	66.4% ± 2.0%
FedAvg	81.6% ± 0.2%	84.2% ± 0.5%	97.1% ± 1.0%	95.7% ± 0.5%	81.6% ± 3.0%	78.2% ± 0.7%	68.3% ± 2.6%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle3 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3% | specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	84.3% ± 1.4%	86.1% ± 1.1%	96.7% ± 0.5%	95.8% ± 1.1%	86.6% ± 2.6%	78.0% ± 2.0%	73.7% ± 3.9%
FedAvg	81.5% ± 0.2%	83.9% ± 0.5%	95.4% ± 0.8%	95.8% ± 1.0%	82.4% ± 3.4%	77.1% ± 0.4%	68.7% ± 2.8%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle4 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3% | specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.9% ± 2.9%	84.6% ± 2.1%	96.4% ± 1.2%	96.1% ± 0.4%	85.9% ± 3.7%	76.4% ± 0.2%	68.1% ± 7.4%
FedAvg	81.5% ± 0.6%	84.0% ± 0.3%	96.2% ± 1.0%	94.5% ± 0.8%	82.6% ± 2.1%	78.3% ± 0.6%	68.1% ± 3.3%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle5 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 5.0%±0.0% / true 6.3%
specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	85.4% ± 1.5%	87.2% ± 0.5%	96.5% ± 1.7%	95.7% ± 0.7%	88.3% ± 2.6%	81.1% ± 0.7%	74.6% ± 7.4%
FedAvg	81.2% ± 0.3%	83.5% ± 0.4%	95.1% ± 1.1%	94.9% ± 1.1%	81.6% ± 3.1%	77.8% ± 0.9%	68.3% ± 2.6%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle6 | n_clients=1 | cycle2=visible 4.6%±0.0% / true 5.8% | cycle3=visible 5.0%±0.0% / true 6.2% | cycle4=visible 5.3%±0.0% / true 6.6% | cycle5=visible 5.1%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3%
specialization=low | specialization_fraction=0.20 | global_visible_support=0.84 | task_jsd_mean=0.003 | task_jsd_median=0.003 | task_jsd_max=0.003 | family_centroid_jsd_mean=0.003


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	84.2% ± 1.5%	86.2% ± 1.0%	97.7% ± 0.7%	93.8% ± 0.8%	88.3% ± 1.2%	79.5% ± 0.8%	71.9% ± 4.5%
FedAvg	82.2% ± 0.7%	84.7% ± 0.3%	97.4% ± 0.3%	96.0% ± 1.1%	83.2% ± 2.3%	78.5% ± 0.9%	68.6% ± 2.8%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




Controlled task-specialization level: Mid | FedAvg diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64

family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
family=specialized_cycle3 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
family=specialized_cycle4 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
family=specialized_cycle5 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
family=specialized_cycle6 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 2.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3%
specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.6% ± 0.7%	84.6% ± 0.7%	96.9% ± 0.3%	95.3% ± 0.9%	87.5% ± 0.4%	77.3% ± 1.9%	65.6% ± 0.6%
FedAvg	78.7% ± 2.2%	81.3% ± 1.4%	94.2% ± 1.3%	92.5% ± 1.6%	79.5% ± 0.1%	74.6% ± 3.1%	65.8% ± 4.6%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3%
specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.3% ± 1.7%	84.2% ± 1.3%	98.2% ± 0.7%	94.8% ± 1.7%	86.7% ± 1.3%	75.0% ± 3.4%	66.4% ± 3.0%
FedAvg	78.8% ± 1.9%	81.4% ± 1.1%	94.8% ± 1.6%	92.9% ± 1.0%	79.0% ± 0.6%	74.7% ± 2.7%	65.7% ± 4.6%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle3 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3% | specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.1% ± 0.2%	84.7% ± 0.1%	97.7% ± 0.8%	96.9% ± 0.8%	90.8% ± 0.3%	74.1% ± 1.7%	64.0% ± 0.4%
FedAvg	78.7% ± 2.1%	81.3% ± 1.4%	93.5% ± 0.9%	92.5% ± 1.6%	80.1% ± 0.7%	74.1% ± 2.8%	65.9% ± 4.7%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle4 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3% | specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	82.3% ± 1.4%	85.7% ± 1.2%	97.7% ± 1.1%	96.2% ± 0.1%	89.8% ± 1.6%	80.9% ± 4.0%	63.8% ± 1.5%
FedAvg	78.7% ± 2.2%	81.3% ± 1.3%	94.8% ± 2.3%	91.9% ± 1.2%	79.7% ± 0.4%	74.3% ± 3.1%	65.7% ± 4.8%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle5 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 3.1%±0.0% / true 6.3% | specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	82.2% ± 0.8%	84.9% ± 0.9%	95.7% ± 1.8%	95.1% ± 0.8%	86.8% ± 3.3%	80.6% ± 1.1%	66.1% ± 1.3%
FedAvg	78.2% ± 2.1%	80.7% ± 1.3%	93.4% ± 1.5%	91.9% ± 1.8%	78.5% ± 0.4%	74.5% ± 3.9%	65.4% ± 4.6%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle6 | n_clients=1 | cycle2=visible 2.9%±0.0% / true 5.8% | cycle3=visible 3.1%±0.0% / true 6.2% | cycle4=visible 3.3%±0.0% / true 6.6% | cycle5=visible 3.2%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3% | specialization=mid | specialization_fraction=0.50 | global_visible_support=0.60 | task_jsd_mean=0.028 | task_jsd_median=0.028 | task_jsd_max=0.029 | family_centroid_jsd_mean=0.028


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	81.3% ± 2.7%	83.7% ± 2.0%	95.6% ± 1.4%	93.8% ± 1.7%	84.2% ± 3.5%	76.6% ± 4.7%	68.5% ± 6.6%
FedAvg	79.3% ± 2.6%	81.9% ± 1.9%	94.6% ± 1.4%	93.3% ± 2.5%	80.0% ± 1.1%	75.2% ± 3.1%	66.1% ± 4.7%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




Controlled task-specialization level: High | FedAvg diagnostics | aggregated / pooled over 5 clients | model=mcwauto_layers6_lr0.001_wd0.0001_do0.1_hd64_egoTrue_bs64

family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=specialized_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=specialized_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=specialized_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
family=specialized_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3%
specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 | family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	74.7% ± 1.4%	78.7% ± 1.0%	91.5% ± 2.6%	89.1% ± 0.9%	80.9% ± 1.9%	68.6% ± 3.9%	58.8% ± 1.7%
FedAvg	76.8% ± 1.8%	80.5% ± 1.9%	95.1% ± 1.3%	92.3% ± 3.1%	81.6% ± 1.4%	74.6% ± 3.1%	59.0% ± 1.0%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle2 | n_clients=1 | cycle2=visible 5.8%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3% | specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 | family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	71.8% ± 2.2%	76.9% ± 0.8%	98.1% ± 0.8%	90.4% ± 2.4%	76.9% ± 4.5%	65.2% ± 5.3%	53.7% ± 4.1%
FedAvg	77.0% ± 1.8%	80.7% ± 1.8%	95.6% ± 0.6%	92.2% ± 3.4%	81.6% ± 1.2%	74.7% ± 2.9%	59.3% ± 1.1%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle3 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 6.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% |
cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 |
family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	69.6% ± 3.7%	76.6% ± 3.0%	96.4% ± 1.1%	94.7% ± 1.9%	86.7% ± 5.4%	52.7% ± 5.2%	52.7% ± 4.1%
FedAvg	76.9% ± 1.3%	80.7% ± 1.4%	94.9% ± 1.8%	93.0% ± 2.3%	82.0% ± 0.5%	74.4% ± 2.1%	59.0% ± 1.0%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle4 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 6.6%±0.0% / true 6.6% |
cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 |
family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	78.0% ± 1.1%	82.8% ± 1.1%	95.4% ± 2.6%	95.6% ± 1.0%	91.0% ± 2.8%	76.7% ± 2.4%	55.2% ± 0.7%
FedAvg	76.7% ± 1.9%	80.5% ± 1.9%	95.2% ± 0.1%	92.3% ± 3.6%	81.3% ± 1.8%	74.7% ± 3.4%	59.0% ± 1.3%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle5 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% |
cycle5=visible 6.4%±0.0% / true 6.4% | cycle6=visible 1.3%±0.0% / true 6.3%
specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 |
family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	75.8% ± 0.9%	78.3% ± 1.5%	87.0% ± 1.3%	87.4% ± 3.9%	77.0% ± 3.6%	79.2% ± 5.6%	60.7% ± 3.5%
FedAvg	76.5% ± 1.9%	80.2% ± 2.0%	94.7% ± 0.9%	91.3% ± 3.3%	81.3% ± 1.5%	74.4% ± 3.6%	59.0% ± 1.0%


Validation loss comparison




Pooled local centralized micro / macro F1




Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)




family=specialized_cycle6 | n_clients=1 | cycle2=visible 1.2%±0.0% / true 5.8% | cycle3=visible 1.2%±0.0% / true 6.2% | cycle4=visible 1.3%±0.0% / true 6.6% | cycle5=visible 1.3%±0.0% / true 6.4% | cycle6=visible 6.3%±0.0% / true 6.3% | specialization=high | specialization_fraction=0.80 | global_visible_support=0.36 | task_jsd_mean=0.162 | task_jsd_median=0.162 | task_jsd_max=0.165 | family_centroid_jsd_mean=0.162


Best logged test performance (FedAvg vs Local centralized)

Algorithm	Micro-F1	Macro Pos-F1	cycle2 pos-F1	cycle3 pos-F1	cycle4 pos-F1	cycle5 pos-F1	cycle6 pos-F1
Local centralized	78.5% ± 3.5%	79.1% ± 2.8%	83.3% ± 8.5%	79.7% ± 5.3%	76.3% ± 4.2%	77.7% ± 6.0%	78.4% ± 6.9%
FedAvg	77.0% ± 2.2%	80.7% ± 2.4%	95.2% ± 3.0%	92.8% ± 2.9%	81.7% ± 2.4%	75.0% ± 3.6%	58.9% ± 1.0%


Validation loss comparison




Pooled local centralized micro / macro F1



Pooled FedAvg micro / macro F1 (pre-aggregation)



Pooled FedAvg micro / macro F1 (post-aggregation)



Delta vs FedAvg (local minus pre / post)

